

# Game Technology

Lecture 11 – 16.01.2015



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

# Preliminary timetable

Lecture No.	Date	Topic
1	17.10.2014	Basic Input & Output
2	24.10.2014	Timing & Basic Game Mechanics
3	31.10.2014	Software Rendering 1
4	07.11.2014	Software Rendering 2
5	14.11.2014	Basic Hardware Rendering
6	21.11.2014	Animations
7	28.11.2014	Physically-based Rendering
8	05.12.2014	Physics 1
9	12.12.2014	Physics 2
10	19.12.2014	Procedural Content Generation
<b>11</b>	<b>16.01.2015</b>	<b>Compression &amp; Streaming</b>
12	23.01.2015	Multiplayer
13	30.01.2015	Audio
14	06.02.2015	Scripting
15	13.02.2015	AI



# Today's Games

---

## Typical hardware requirements

- 8 GiB RAM
- 2 GiB Video-RAM
- 50 GiB on disk

## All SNES games ever (including all language versions)

- ~3000 games
- ~4.5 GiB

## One uncompressed texture

- $4096 \times 4096 \times 4 \text{ Bytes} = 67108864 \text{ Bytes} = 64 \text{ MiB}$
- $2 \text{ GiB} / 64 \text{ MiB} = 32$
- Physically based rendering – typically 4 textures

## Killzone 4 CPU data



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

Sound	553 MB
Havok Scratch	350 MB
Game Heap	318 MB
Various Assets, Entities, etc.	143 MB
Animation	75 MB
Executable + Stack	74 MB
LUA Script	6 MB
Particle Buffer	6 MB
AI Data	6 MB
Physics Meshes	5 MB
<b>Total</b>	<b>1,536 MB</b>

# Killzone 4 GPU data

Non-Steaming Textures	1,321 MB
Render Targets	800 MB
Streaming Pool (1.6 GB of streaming data)	572 MB
Meshes	315 MB
CUE Heap (49x)	32 MB
ES-GS Buffer	16 MB
GS-VS Buffer	16 MB
<b>Total</b>	<b>3,072 MB</b>



# PNG and JPEG

---

## PNG

- Lossless
- Compression highly dependent on image content

## JPEG

- Lossy
- Generally strong compression

## Both

- Slow decompression
  - Can slow down loading times
- Not possible to access a single pixel while compressed
  - Not usable for image computations  
aka not usable as a texture format



# Texture Compression

---

## Many different formats

- S3TC, PVRTC, ASTC,...
- Has to be supported by GPU and Graphics API
- Of course much of it is patented and hard to standardize

## Design goals

- High compression
- Low visual degradation
- Efficient single pixel access
  - Constant size of a pixel or a pixel block

# Possible compression strategies

---

**Less than 8 bits per color might be ok**

**The eye's color resolution is less than its intensity resolution**

**Neighboring pixels likely have similar colors**

# Example 1

---

## ETC

- Ericsson Texture Compression

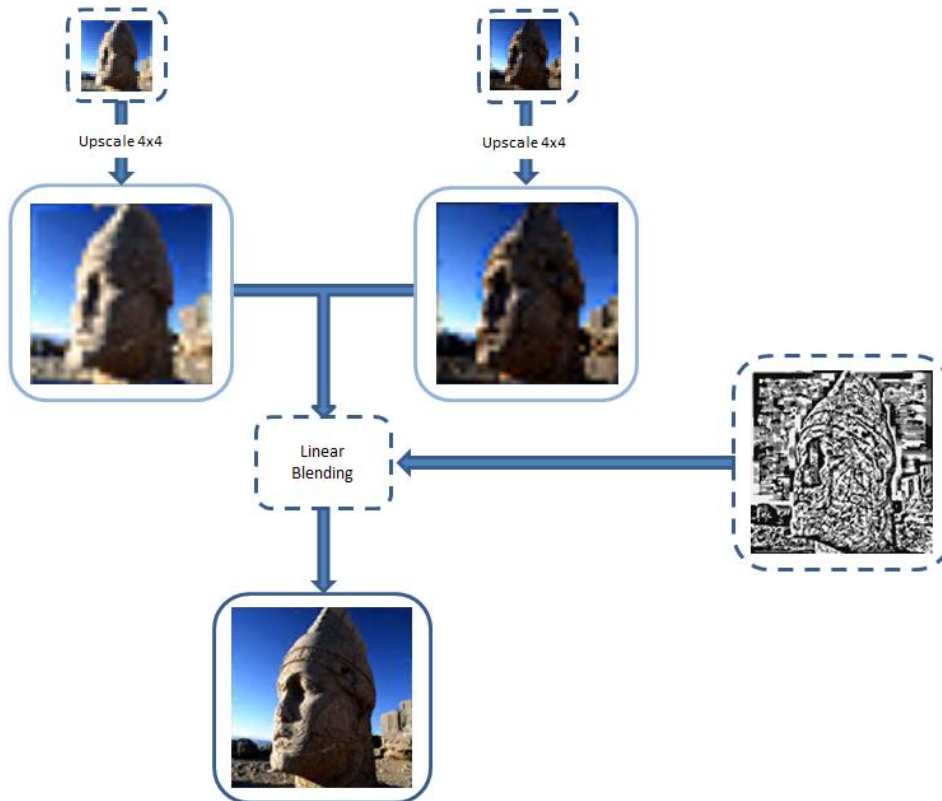
### **Compresses 4x4 pixel blocks to 64 bits**

- Split into two 2x4 groups
- Each gets a 12 bit base color plus 3 bit brightness range selection
- Each pixel gets a 2 bit offset value

# Example 2

## PVRTC

- PowerVR Texture Compression



# Normal Maps,...

## Compression for images might not be optimal for other textures

- But it might just work
- Swizzling channels can help

## 3Dc

- $x^2 + y^2 + z^2 = 1$ 
  - $z^2 = 1 - x^2 - y^2$
  - One value can be omitted
- Plus block compression

# Manual Compression

---

**Let the artists do the job**


**Repeat images over and over**

- Nobody might notice it when you do it cleverly

# Tilemaps/Tilesets




TECHNISCHE  
UNIVERSITÄT  
DARMSTADT



# Tile Editors



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT





# Pitfall: The Mayan Adventure



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT



# Warcraft 3



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT



## Bilinear Filtering

- Would have to use texels from two tiles at tile boundaries
- Complicated
- Expensive
- Rarely used

# Multitexturing



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT





# Multitexturing



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT






# Multitexturing



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT



## Good lighting can hide a lack of details



## Performance

- More textures, less performance
- Precalculating which polys actually use more textures can help

## Needs good tool support

- Scary communication with artists



## Coarse Streaming

- Load and replace complete assets

## Fine Grained Streaming

- Load and show/play a single asset bit by bit

## Similar to level of detail systems

- Load big textures for near objects
- Kick out big textures for far away objects
- Maybe blend texture changes in and out





# Problems

---

## **Disks are slow and unreliable**

- No timing guarantees at all
- Load textures in a second thread,  
always have an emergency strategy ready  
(keep super low resolution textures of everything in RAM)

## **Changing textures at runtime is problematic**

- Driver might decide to convert the texture
- Easier on console
- Probably easier with Direct3D 12

# Fine grained texture streaming



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT





# MegaTextures

## Really huge textures

- Rage supports textures of up to 128000x128000
  - That's ~60 GiB

## Compression

- Texture is highly compressed on disk
  - Using lossy JPEG like compression

## One texture for everything

- Complete world in one texture
- No restrictions for artists
  - But toolsets provide classical multitexturing tricks
  - Artists don't manually paint 128000x128000 pixels



# MegaTextures

## Geometry is split up in tiles

- Engine determines screen size of visible tiles
- Loads texture parts in varying sizes to optimize current view






# Geometry



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT





**Not widely used**

**No hardware support**


**Special strategies for animations**

- Like skeletal animations, which are tiny

# Manual Compression



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT




# Height Maps



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

## Just Y instead of X/Y/Z



# Normal Maps

---

**Remove super detailed geometry**

**Replace with normal maps**

- Which is a form of compression by itself
- Plus normal can be further compressed

# Coarse Geometry Streaming

---

## Same strategies as for textures

- Could be directly plugged into a level of detail system

# Fine grained geometry streaming

---

**To be done**

## **mp3 and similar compressed formats**

- Nothing special – at least not anymore

## **Coarse streaming for sound effects**

- Easy
  - Sound effects are short
  - Sound effects don't stay on screen
  - Sound effects can stay in CPU RAM

## **Fine grained streaming for music and maybe speech**

- Even mp3 players do it

## 32 bit floats

- “total precision is 24 bits (equivalent to  $\log_{10}(2^{24}) \approx 7.225$  decimal digits)”
  - Can be a little tight for big worlds

## Use 64 bit floats for positions

- Hard to integrate 32 bit physics engines

## Split and Shift the world

- Split the world
- Shift the closest parts to a position nearer at the camera



## Sampling

- Samples at random intervals
- Does not modify code

## Instrumentation

- Adds sampling code to binary

**CPU integrated circuitry that measures  
certain performance characteristics**

- Like number of cache misses,...

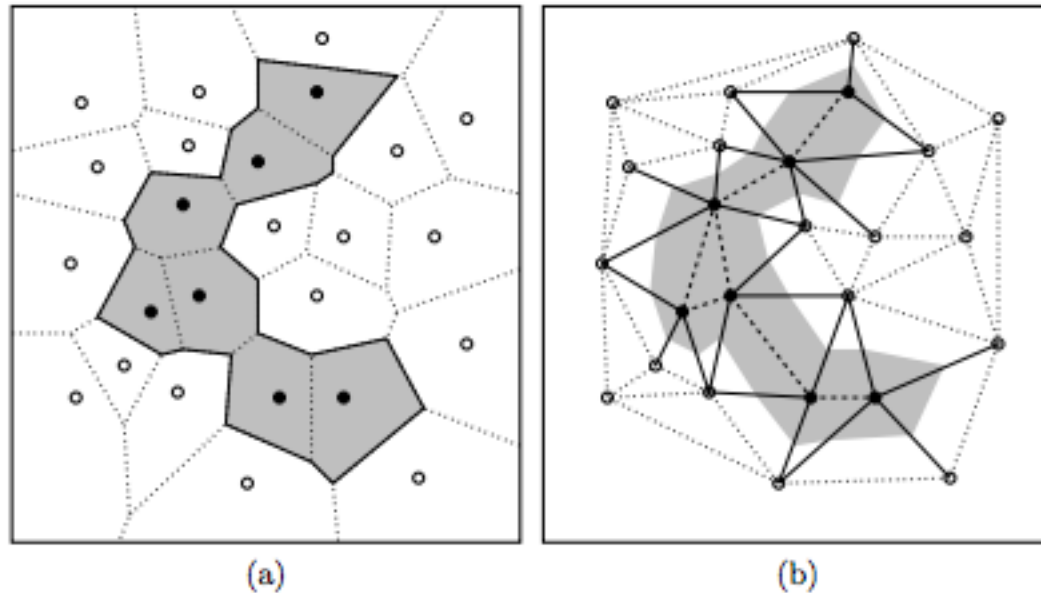
**Can be read by CPU specific profiling tools**

## Task 2.1

**We are trying to find a „non-convex hull“ of the points of the cities**  
**A regular convex hull would not approximate our mental image of a**  
**„border“ between two regions**

## Task 2.1

**Use the cells that contain the cities and visualize those areas**  
**Maybe blur or soften the edges to make them rounder**  
**What about „islands“?**



## Task 2.1

---

**See Paper „What Is the Region Occupied by a Set of Points?” by  
Antony Galton and Matt Duckham for more info on the problem of  
“non-convex hulls”  
(<http://www.geosensor.net/papers/galton06.GISCIENCE.pdf>)**


## Task 2.2

**Easier to see in 1D**

**An edge is a sudden change of function value**

**1D kernel:**

- -1 2 -1





## Task 2.2

**If all three pixels are constant:**

- $p_1 = p_2 = p_3 = c$
- $\text{Result} = -1c + 2c - 1c$
- $= 2c - 2c = 0$

**If the left pixel is different**

- $p_1 = c_1, p_2 = p_3 = c_2$
- $\text{Result} = -1c_1 + 2c_2 - c_2$
- $= c_2 - c_1$

**→ Constant areas of the image become black, the edges remain**

## Task 2.3

**With larger sigma, the curve doesn't reach as high, becomes more stretched out**

**The area under the curve stays the same**

**→ The curve becomes wider and is therefore visibly cut off and if we use the coefficients without normalizing, they will not add up to 1**

**→ Without normalizing, we will change the brightness of the image**

